

# Generalization Bounds

## Theoretical Foundations of Deep Learning

Matteo Mazzarelli

December 17, 2024



# What is Machine Learning?

- ▶ Machine learning is the process of learning from data to make predictions or decisions. [1]

# Supervised Learning

- ▶ We focus on supervised learning where the data consists of input-output pairs, called features ( $x_i$ ) and labels ( $y_i$ ). [1-3]
- ▶ The goal is to infer  $y_i$  from  $x_i$ . [4]

# The Learning Problem

- ▶ We have a dataset of observations  
 $S = \{(x_1, y_1), \dots, (x_m, y_m)\}$ . [4]
- ▶ We wish to learn how to infer the value of  $y_i$  given  $x_i$ . [4]

# Statistical Model

- ▶ We assume the values of  $(x_i, y_i)$  in the dataset are a random sample from a larger population. [5]
- ▶ The values of  $x_i$  and  $y_i$  are realizations of two random variables  $X$  and  $Y$  with probability distributions  $P_X$  and  $P_Y$  respectively. [5]

# Joint Distribution

- ▶ There is a relation between the features and the labels. [6]
- ▶ The value of  $Y$  is conditioned on the value of  $X$ . [6]
- ▶ This is expressed by the conditional probability  $P(Y|X)$ . [6]
- ▶ We can compress  $P_X$  and  $P_Y$  into a single joint distribution  $P(X, Y) = P(X)P(Y|X)$ . [6]

# Target Function

- ▶ The **target function** is  $f(X) = \mathbb{E}[Y|X]$ . [7]
- ▶ It represents the expected value of the label  $Y$  given the features  $X$ . [7]
- ▶ It becomes the target of the machine learning process. [7]
- ▶ The goal is to estimate this function  $f$ . [7]

# Hypothesis and Hypothesis Space

- ▶ A **hypothesis**, denoted by  $h$ , attempts to estimate the target function  $f$ . [7]
- ▶ The **hypothesis space**, denoted by  $\mathcal{H}$ , is the set of possible functions considered for  $h$ . [7, 8]



# Empirical Risk

- ▶ The **empirical risk**, or training error,  $R_{emp}(h)$ , is the average loss of a hypothesis  $h$  on the training data. [9]
- ▶ It can be calculated using a loss function that quantifies the difference between the predicted and actual labels. [9]

# Overfitting

- ▶ Simply minimizing empirical risk can lead to **overfitting**. [10]
- ▶ An overfit model performs well on the training data, but poorly on unseen data. [10, 11]
- ▶ This occurs when the model learns the specific details of the training data instead of the underlying patterns. [10]

# Generalization Error

- ▶ The **generalization error** (risk),  $R(h)$ , measures how well the hypothesis  $h$  performs on unseen data. [12]
- ▶ It's the expected value of the loss over the entire joint distribution  $P(X, Y)$ . [12]

# Generalization Gap

- ▶ The **generalization gap** is the difference between the training error and the generalization error. [13]
- ▶ It quantifies how well the performance on the training data generalizes to unseen data. [13]

# Motivation

- ▶ **Generalization bounds** provide guarantees that the learned hypothesis will perform well on unseen data. [2, 14, 15]
- ▶ They relate the generalization error to quantities we can observe or control, such as empirical risk, hypothesis space complexity, and dataset size. [16]

# Hoeffding's Inequality

- For a single hypothesis  $h$ , Hoeffding's inequality bounds the difference between the empirical risk and the generalization error. [17, 18]

# Limitations of Hoeffding

- ▶ Hoeffding's inequality doesn't directly apply to the entire hypothesis space. [8]
- ▶ It only considers the boundedness of the functions, not their variance. [19]
- ▶ The union bound, used to extend Hoeffding to multiple hypotheses, assumes all hypotheses are independent, which is not generally true. [20, 21]

# The Union Bound

- ▶ The **union bound** extends the probability bounds to the entire hypothesis space. [22]
- ▶ It states that the probability of at least one hypothesis having a large generalization gap is at most the sum of the probabilities of each individual hypothesis having a large gap. [22]



# The Growth Function

- ▶ The **growth function** quantifies the expressiveness of the hypothesis space. [23]
- ▶ It's the maximum number of ways the hypothesis space can label a dataset of a given size. [23]

# VC Dimension

- ▶ The **VC dimension** is the largest dataset size that the hypothesis space can **shatter**. [24, 25]
- ▶ Shattering means the hypothesis space can produce all possible labelings for the dataset. [24]
- ▶ It's a measure of the complexity of the hypothesis space. [26]

# VC Generalization Bound

- ▶ The **VC generalization bound** relates the generalization error to the empirical risk, VC dimension, and dataset size. [27]
- ▶ It shows that the generalization error can be bounded by the empirical risk plus a term that depends on the VC dimension and dataset size. [28]

## Distribution-Based Bounds

- ▶ VC dimension is **distribution-free**, meaning it doesn't consider the data distribution. [29]
- ▶ Tighter bounds can be obtained by considering the data distribution. [29, 30]
- ▶ **Support Vector Machines (SVMs)** exemplify this by maximizing the margin between classes, which leads to a lower VC dimension and better generalization. [30]

## Other Capacity Measures

- ▶ **Covering numbers** measure the size of the hypothesis space using a metric based on the difference in predictions on the training data. [31]
- ▶ **Rademacher complexity** measures how well the hypothesis space can fit random noise. [32, 33]
- ▶ These measures can be used to derive generalization bounds. [32, 34]

# The General Form of Generalization Bounds

- ▶ The general form of generalization bounds is:  
$$R(h) \leq R_{emp}(h) + C(|\mathcal{H}|, N, \delta). \quad [35]$$
- ▶  $C(|\mathcal{H}|, N, \delta)$  represents a complexity term that depends on the hypothesis space complexity, dataset size, and the desired confidence level. [35]
- ▶ It highlights the trade-off between minimizing the training error and controlling the model's complexity. [35]

# Relevance to Other Topics

## Rates of Convergence

- ▶ **Rates of convergence** quantify how fast the generalization error decreases with increasing sample size. [36]
- ▶ They are closely linked to generalization bounds, as the bounds often provide insights into the rate of convergence. [37]

## PAC-Bayes

- ▶ **PAC-Bayes** offers a Bayesian approach to deriving generalization bounds. [38]

# Key Takeaways

- ▶ **Generalization bounds** are crucial for understanding and controlling the performance of machine learning models. [15]
- ▶ They guarantee the learned hypothesis will perform well on unseen data. [15]