

Generalization Bounds

Theoretical Foundations of Deep Learning

Matteo Mazzarelli

December 17, 2024



Motivation

- ▶ **Core Challenge:** How can a model learned from *limited training data* perform well on *unseen data*?
- ▶ Generalization lies at the heart of the machine learning process.
- ▶ A poorly generalized model risks:
 - ▶ **Overfitting:** Performing well on training data but poorly on unseen data.
 - ▶ **Underfitting:** Failing to capture the underlying patterns of the data.

The Learning Problem

▶ Supervised Learning:

- ▶ Goal: Learn a function ($f: X \rightarrow Y$) mapping inputs (X) to outputs (Y) based on labeled training data.

▶ Key Question: Can the learned function perform well on unseen data?

▶ Generalization:

- ▶ Ability of a model to extend its learning beyond the training data.
- ▶ **Central Problem** in machine learning: balancing *empirical performance* with *future predictions*.

Why Theory Matters

- ▶ **Significance of Theory:**
 - ▶ Guides **algorithm design** by providing a foundation for developing new methods.
 - ▶ Allows **performance analysis** to identify the strengths and weaknesses of algorithms.
 - ▶ Reveals **limitations** of learning systems, helping us understand their boundaries.
- ▶ **Theoretical Understanding:**
 - ▶ Bridges the gap between empirical performance and guarantees on future behavior.

Introducing Generalization Bounds

► What Are Generalization Bounds?

- Theoretical tools offering guarantees about a model's performance on unseen data.
- Relate:
 - **Generalization Error:** How well the model generalizes.
 - **Empirical Risk:** Performance observed on training data.
 - **Model Complexity:** How expressive the model is.

► Purpose:

- Provide insights into the trade-offs between model accuracy, complexity, and training data size.