

Generalization Bounds

Theoretical Foundations of Deep Learning

Matteo Mazzarelli

December 17, 2024



Introduction

Motivation

- ▶ **Core Question:** How can models trained on limited data perform reliably on unseen scenarios?
- ▶ **Generalization** is a fundamental goal in machine learning: ensuring models extend their learned patterns to new, unseen data.
- ▶ A poorly generalized model risks:
 - ▶ **Overfitting:** Performing well on training data but poorly on unseen data.
 - ▶ **Underfitting:** Failing to capture the underlying patterns of the data.

The Learning Problem

- ▶ **Supervised Learning:**
 - ▶ Goal: Learn a function $f : X \rightarrow Y$ mapping inputs X to outputs Y based on labeled training data.
- ▶ **Key Question:** Can the learned function perform well on unseen data?
- ▶ **Generalization:**
 - ▶ Ability of a model to extend its learning beyond the training data.
 - ▶ **Central Problem** in machine learning: balancing *empirical performance* with *future predictions*.

Overfitting

Demonstrating Overfitting

- ▶ **Objective:**
 - ▶ Show how increasing model complexity (polynomial degree) leads to overfitting.
- ▶ **Dataset:**
 - ▶ Using the scikit-learn **Diabetes** dataset with a single feature (BMI) and a quantitative response variable indicating disease progression (Target)^[1].
- ▶ **Approach:**
 1. Fit polynomial regression models of varying degrees.
 2. Visualize polynomial fits on the training data.
 3. Examine the fits' residuals to see how errors behave.
 4. Plot training vs. test errors to highlight overfitting.

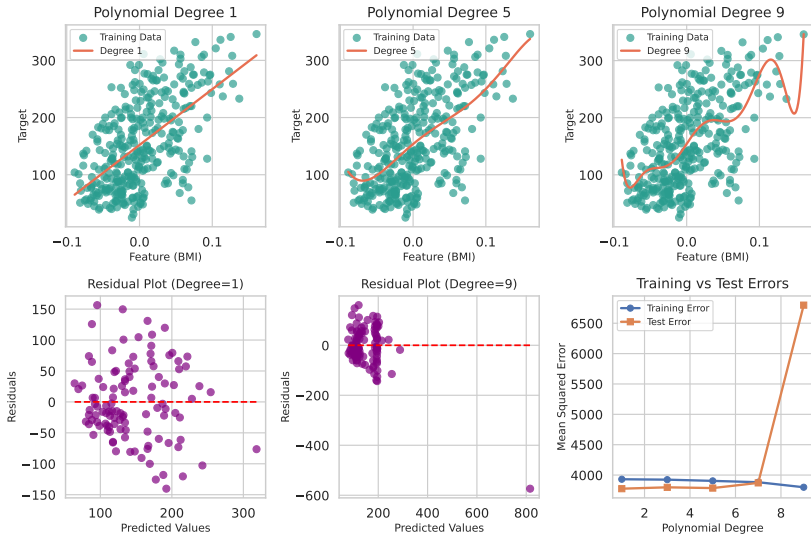


Figure 1: Overfitting Phenomenon in Polynomial Regression

Double Descent

- Modern machine learning introduces a fascinating twist: **Double Descent**, where increasing model complexity can lead to improved generalization after an initial overfitting phase.



Figure 2: Double Descent phenomenon in a Residual Neural Network^[2]

Classical Bounds

Introducing Generalization Bounds

- ▶ **What Are Generalization Bounds?**
 - ▶ Theoretical tools offering guarantees about a model's performance on unseen data.
 - ▶ Relate:
 - ▶ **Generalization Error**: How well the model performs on unseen data.
 - ▶ **Empirical Risk**: Performance observed on training data.
 - ▶ **Model Complexity**: How expressive the model is.

► **Purpose:**

- Provide insights into the trade-offs between:
 - **Model Accuracy:** How well the model captures the data patterns.
 - **Model Complexity:** The expressiveness of the model and its ability to fit intricate patterns.
 - **Training Data Size:** How much data is required to achieve reliable generalization.

Hoeffding's Inequality: A Starting Point

- ▶ **What is Hoeffding's Inequality?**
 - ▶ A fundamental result in probability theory used to bound the difference between the **empirical risk** and the **generalization error** for a fixed hypothesis.
 - ▶ Provides a way to measure how closely a model's performance on training data reflects its performance on unseen data.

Mathematical Formulation of Hoeffding's Inequality

► Hoeffding's Inequality:

$$P(|R(h) - R_{\text{emp}}(h)| > \varepsilon) \leq 2 \exp(-2m\varepsilon^2)$$

- $R(h)$: Generalization error (true performance on unseen data).
- $R_{\text{emp}}(h)$: Empirical risk (error on training data).
- ε : A small positive value (tolerance).
- m : Size of the dataset.

Key Insights

- ▶ The probability that the generalization error $R(h)$ deviates significantly from the empirical risk $R_{\text{emp}}(h)$ decreases **exponentially** with:
 - ▶ Larger dataset size m .
 - ▶ Smaller tolerance ε .

Rates of Convergence

► What Are Rates of Convergence?

- Quantify how quickly the generalization error approaches the empirical risk as the dataset size m grows.
- Provide a guideline for determining the dataset size needed to achieve a desired level of generalization.
- In Hoeffding's inequality:

$$P(|R(h) - R_{\text{emp}}(h)| > \varepsilon) \leq 2 \exp(-2m\varepsilon^2)$$

- The **exponential term** $\exp(-2m\varepsilon^2)$ shows that the convergence is faster with larger datasets.

► Key Factors:

- **Dataset Size (m)**: Larger datasets reduce the gap between $R(h)$ and $R_{\text{emp}}(h)$ more quickly.
- **Tolerance (ε)**: Smaller tolerances require larger datasets for the same level of confidence.

Interpretation of Hoeffding's Inequality

► What Does It Mean?

- As the dataset size (m) increases, the empirical risk becomes a more reliable indicator of the generalization error.
- For a fixed hypothesis, we can be confident that the performance observed on training data is close to what can be expected on unseen data.
- The **rate of convergence** shows how quickly this reliability improves as m grows.

Key Insights

- ▶ Hoeffding's inequality gives a **quantitative guarantee** about the relationship between training performance and unseen data performance.
- ▶ Understanding convergence rates helps in planning how much data is needed for robust generalization.

Limitations of Hoeffding's Inequality

- ▶ **Beyond Fixed Hypotheses:**
 - ▶ Hoeffding's inequality assumes a single, fixed hypothesis.
 - ▶ In practice, machine learning involves selecting the best hypothesis from a **large hypothesis class** \mathcal{H} , increasing the risk of overfitting.
- ▶ **Need for Complexity-Aware Bounds:**
 - ▶ Simple bounds like Hoeffding's do not consider the complexity of the hypothesis class, which influences generalization.

The Union Bound

- ▶ **What is the Union Bound?**
 - ▶ A probability tool used to extend bounds like Hoeffding's inequality to apply across an entire hypothesis space \mathcal{H} .
 - ▶ Helps estimate the probability that **at least one hypothesis** in \mathcal{H} has a large generalization gap.
- ▶ **Key Idea:**
 - ▶ Instead of considering a single fixed hypothesis, the Union Bound aggregates the probabilities of generalization gaps over all hypotheses in \mathcal{H} .

Formalization of The Union Bound

► Mathematical Expression:

$$P\left(\sup_{h \in \mathcal{H}} |R(h) - R_{\text{emp}}(h)| > \epsilon\right) \leq \sum_{h \in \mathcal{H}} P(|R(h) - R_{\text{emp}}(h)| > \epsilon)$$

- $\sup_{h \in \mathcal{H}}$: The supremum ensures we account for the worst-case scenario across all hypotheses.
- $P(|R(h) - R_{\text{emp}}(h)| > \epsilon)$: The probability of a significant generalization gap for each hypothesis.
- **How It Works:**
 - By summing up the probabilities for all hypotheses, the Union Bound provides a way to analyze the worst-case scenario over the hypothesis space.

Implications of The Union Bound

- ▶ **Impact of Hypothesis Space Size:**
 - ▶ The bound depends directly on the **size of the hypothesis space** $|\mathcal{H}|$.
 - ▶ Larger hypothesis spaces increase the sum, making the bound looser.
- ▶ **Takeaway:**
 - ▶ The Union Bound highlights a trade-off:
 - ▶ **Small hypothesis space:** Tighter bounds, but limited model capacity.
 - ▶ **Large hypothesis space:** Higher capacity, but risk of overfitting and looser bounds.

Transition to Advanced Bounds

- ▶ **Connection to Practical Learning:**
 - ▶ In practice, hypothesis spaces are often infinite or too large to enumerate explicitly. This motivates the need for alternative ways to measure hypothesis complexity.
- ▶ **From Simple to Sophisticated:**
 - ▶ The Union Bound provides a conceptual basis for understanding how hypothesis space size affects generalization.
 - ▶ Next, we delve into **complexity measures** that allow us to extend generalization bounds to more practical, infinite hypothesis spaces.

Advanced Bounds

Motivation for Advanced Bounds

- ▶ **Advanced bounds** address a variety of the limitations we have outlined by incorporating:
 - ▶ **VC Dimension:** A measure of the capacity or expressiveness of a hypothesis class. Higher VC dimensions indicate more complex models, which may require more data to generalize well.
 - ▶ **Rademacher Complexity:** A data-dependent measure of how well a hypothesis class can fit random noise in the training data. It captures both the hypothesis class and the specifics of the data distribution.

► Extending Convergence Rates:

- Advanced bounds refine the rates of convergence by linking the generalization error to:
 - The size of the dataset m .
 - The complexity of the hypothesis class (e.g., **VC dimension** or **Rademacher complexity**).
- For example, the generalization error is often bounded as:

$$R(h) - R_{\text{emp}}(h) \leq \mathcal{O} \left(\sqrt{\frac{\text{Complexity}(\mathcal{H})}{m}} \right)$$

- Larger datasets m reduce error, but higher complexity increases the required data for a desired level of generalization.

► Practical Implications:

- These bounds provide actionable insights for balancing model complexity and dataset size.

Vapnik-Chervonenkis (VC) Theory

► Growth Function

- The **Growth Function** is a measure of the expressiveness of a hypothesis space \mathcal{H} .
- **Definition:**
 - The growth function, $\Pi_{\mathcal{H}}(m)$, is the maximum number of distinct ways a hypothesis space can label m data points.
- **Key Idea:**
 - A more expressive hypothesis space can label datasets in a greater number of ways, indicating higher complexity.
- **Growth Behavior:**
 - For small m , $\Pi_{\mathcal{H}}(m) = 2^m$.
 - For larger m , the growth may be limited by the structure of \mathcal{H} .

▶ VC Dimension

▶ The **VC Dimension** is a scalar value that quantifies the capacity of a hypothesis space \mathcal{H} .

▶ Definition:

▶ The VC dimension d_{VC} is the size of the largest dataset that can be **shattered** by \mathcal{H} .

▶ Shattering:

▶ A dataset is shattered if every possible labeling of the dataset can be perfectly captured by hypotheses in \mathcal{H} .

▶ Examples:

▶ A linear classifier in 2D space has a VC dimension of 3 (it can shatter any 3 points, but not all configurations of 4 points).

VC Generalization Bound

► What is the VC Generalization Bound?

- A theoretical result that connects the **generalization error** with the **empirical risk**, the **VC dimension**, and the size of the dataset.
- **Mathematical Formulation:**

$$R(h) \leq R_{\text{emp}}(h) + \sqrt{\frac{8d_{\text{VC}} \left(\ln \left(\frac{2m}{d_{\text{VC}}} \right) + 1 \right) + 8 \ln \left(\frac{4}{\delta} \right)}{m}}$$

- $R(h)$: Generalization error.
- $R_{\text{emp}}(h)$: Empirical risk.
- d_{VC} : VC dimension.
- m : Dataset size.
- δ : Confidence level ($1 - \delta$ is the probability that the bound holds).

Key Insights

- ▶ As d_{VC} increases (more complex hypothesis space):
 - ▶ The bound becomes looser, reflecting a higher risk of overfitting.
- ▶ As m increases (larger dataset size):
 - ▶ The bound tightens, improving generalization guarantees.

Summing Up VC Theory

- ▶ **Expressiveness vs. Generalization:**
 - ▶ The VC dimension captures the **expressiveness** of a hypothesis space:
 - ▶ Higher d_{VC} : More complex, more expressive.
 - ▶ A balance is required to avoid overfitting (high complexity) or underfitting (low complexity).
- ▶ **Implications for Learning:**
 - ▶ The VC dimension helps understand:
 - ▶ Why simpler models often generalize better.
 - ▶ Why increasing data size improves generalization, especially for complex models.
- ▶ **Foundation for Algorithm Design:**
 - ▶ VC theory guides the development of learning algorithms by quantifying the trade-offs between hypothesis complexity, data size, and generalization performance.

Distribution-Based Bounds

- ▶ **From General Bounds to Data-Driven Insights:**
 - ▶ Generalization bounds like Hoeffding's inequality and VC bounds rely on worst-case scenarios.
- ▶ **Distribution-Based Bounds:**
 - ▶ Leverage specific properties of the data distribution to achieve **tighter bounds**.
 - ▶ Exploit **data structure** to understand how well a model generalizes in practice.

Example: Support Vector Machines (SVMs)

- ▶ **SVMs and Margin-Based Bounds:**
 - ▶ Support Vector Machines (SVMs) introduce the concept of a **margin**, the distance between the decision boundary and the nearest data points.
 - ▶ **Intuition:**
 - ▶ A larger margin indicates better separation between classes, leading to better generalization.
 - ▶ **Margin-Based Generalization Bounds:**
 - ▶ Generalization error decreases as the margin increases, even for infinite hypothesis spaces.

Alternative Capacity Measures

- ▶ **Why Explore Alternative Measures?**
 - ▶ VC dimension assumes worst-case datasets, often leading to overly conservative bounds.
 - ▶ Alternative measures provide more nuanced insights into hypothesis space complexity, especially for modern machine learning models like neural networks.

► **Examples of Alternative Measures**

1. **Covering Numbers:** The minimum number of small “balls” needed to cover the hypothesis space under a certain metric. Smaller covering numbers indicate a simpler hypothesis space, leading to tighter generalization bounds.
2. **Rademacher Complexity:** Measures the ability of a hypothesis class to fit random noise. A lower Rademacher complexity indicates that the hypothesis space is less prone to overfitting.

► **Next Steps:**

- We want to explore how these theoretical tools are applied to modern machine learning methods.

Key Insights

- ▶ These measures allow for tighter, data-adaptive generalization bounds, particularly useful for complex or large-scale models.
- ▶ There's no one-size-fits-all measure. The choice of capacity measure depends on:
 - ▶ The hypothesis space.
 - ▶ The structure of the data.
 - ▶ The learning algorithm.

Conclusion

References

1. Pedregosa F., Varoquaux G., & et al. (2011). *Scikit-learn: Machine learning in python, diabetes dataset*. https://scikit-learn.org/1.5/modules/generate_d/sklearn.datasets.load_diabetes.html
2. Nakkiran P., Kaplun G., & et al. (2019). *Deep double descent: Where bigger models and more data hurt*. <https://arxiv.org/abs/1912.02292>
3. Bousquet O., Boucheron S., & Lugosi G. (2003). Introduction to statistical learning theory. *Advanced Lectures on Machine Learning*.
4. Samir M. (2016). *A gentle introduction to statistical learning theory*. <https://mostafa-samir.github.io/ml-theory-pt2/>.
5. Vapnik V. N. (1995). *The nature of statistical learning theory*. Springer.
6. Mohri M., Rostamizadeh A., & Talwalkar A. (2012). *Foundations of machine learning*. MIT Press.